


# TensorFlow Eager execution on Backend.AI

: Using TensorFlow Nightly build

Jeongkyu Shin

Nov. 21, 2017





Backend.AI Open-Source PaaS	01
TensorFlow Eager Execution	02
Playing Eager Execution on Backend.AI	03
Closing	04

# Backend.AI OpenSource PaaS



# The era of Open-source ML frameworks





## Cloud

Pay-as-you-go

PaaS for research, deep  
-learning model  
training and ultra-  
convenient coding  
education environment.

## Ground

Bring Your  
Own Hardware

Open-source edition for  
deploying / developing  
your own Backend.AI  
Server Farm.

## Garden

Showcases

Documents, forum,  
showcases of  
Backend.AI platform.

---

## Backend.AI

---

[codeonweb.com](https://codeonweb.com)



## Easy

- Jupyter, VS Code, Atom, IntelliJ plugins
- Only need to set-up API key!

## Fast

- Combine container and GPU technology
- Make sure that the session takes less than a second

## Cheap

- Precise measurement of resource usage up to millisecond / KiB
- (Work In Progress!)

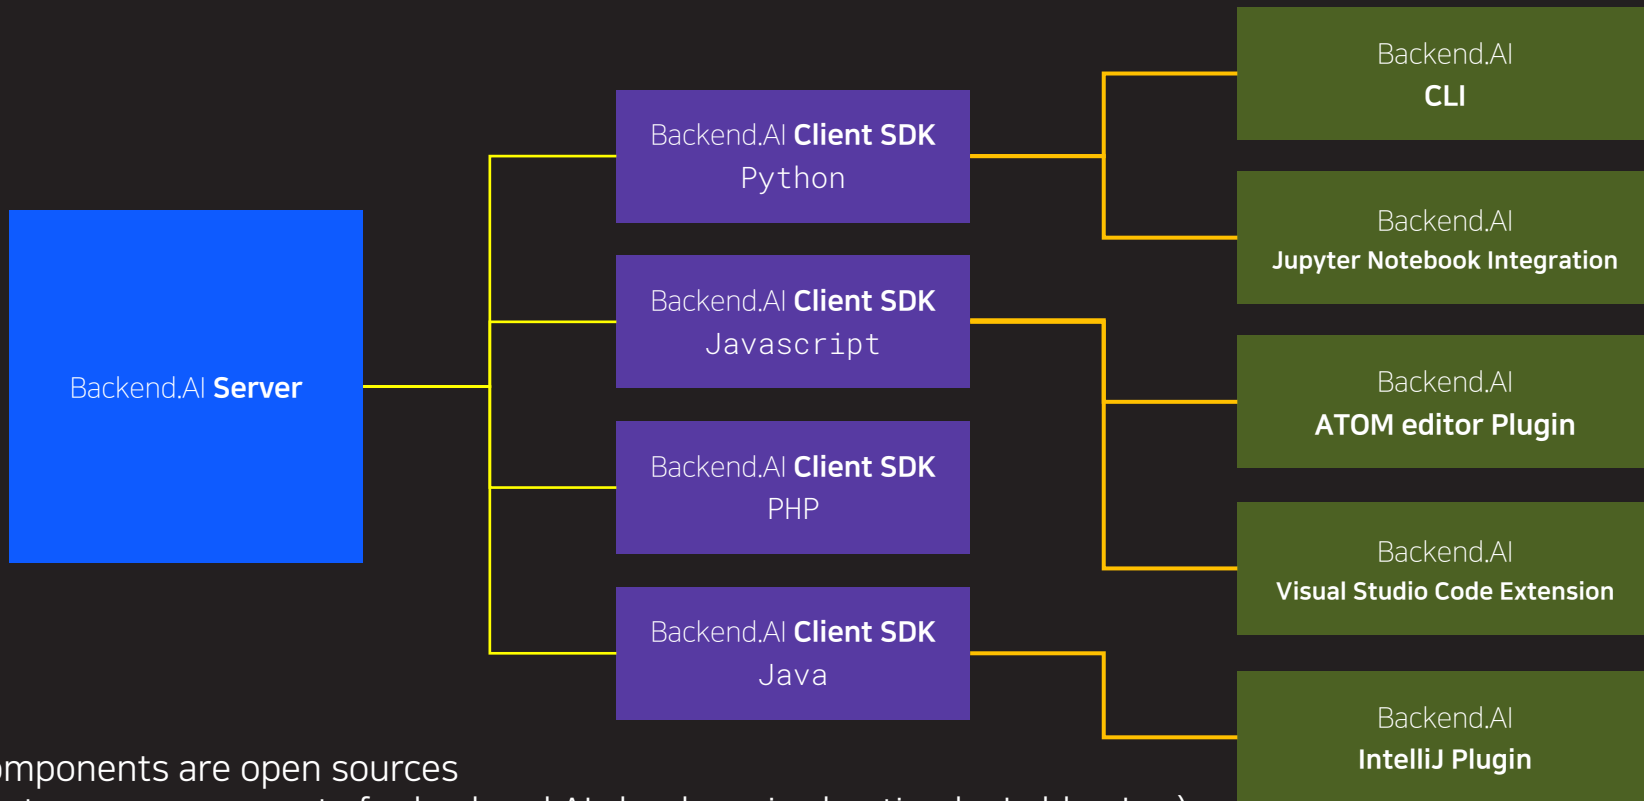
## Together

- Providing a language · version-specific virtual environment using containers
- Syscall sandboxing + enhanced Docker resource limit

## Everywhere

- Open-source version ([github.com/lablup/backend.ai](https://github.com/lablup/backend.ai))
- Cloud service ([cloud.backend.ai](https://cloud.backend.ai))

# Backend.AI Structure



All components are open sources  
(except some components for backend.AI cloud service hosting by Lablup Inc.)

# Backend.AI Core features



- Instant access
  - Available right after issuing API key
  - Creates a virtual environment immediately upon user request
- Handle various requirements
  - Supports all major programming languages and runtimes  
Python, R, Julia, Octave, PHP, Go, C/C++, Java, NodeJS, Lua, Haskell, Rust
  - Supports multiple versions of the same machine learning library  
TensorFlow, Caffe, PyTorch, Keras
- Developer-friendly framework
  - Integrated with familiar user experiences (Editor, IDE, Web notebooks)  
Jupyter, VS Code, Atom Editor, IntelliJ<sup>beta</sup>
  - Provides `$ backend.ai run` CLI, Cloud Interpreter·Compiler
  - HTTP-based public API (REST/GraphQL) and language-specific SDK  
Python, Javascript (Node.js), PHP<sup>beta</sup>





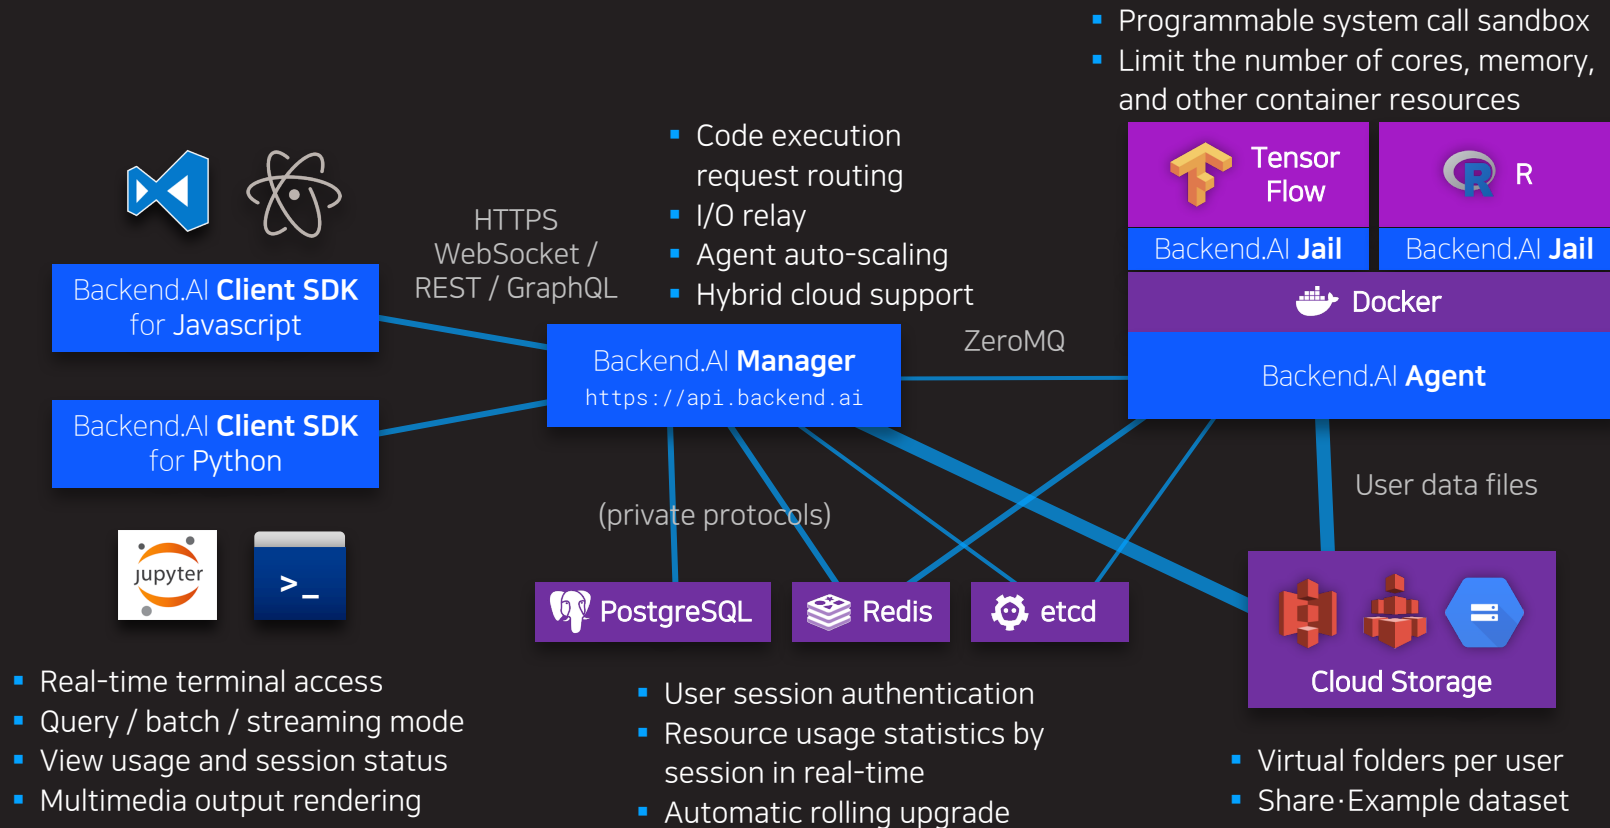
- High performance
  - Asynchronous-based Low latency & high density container pooling (Python asyncio + Docker)
  - GPU-acceleration support for multi-tenant environments (nvidia-docker)
- Security for multi-user / enterprise environments
  - Dynamic sandboxing: programmable syscall filter & rewriting
  - Enforcing resource constraints for legacy apps in Docker
- Hyper-scaling
  - On-premise private cluster
  - Hybrid cloud (on-premise + public cloud)
  - Public cloud (AWS + MS Azure + Google Cloud combination)
  - Route request to specific cloud, based on calculation type, load, or settings

# Language & Frameworks



System	Languages	Frameworks	Developed Framework
Backend.AI <b>Ground</b>	Python 3.6	aiodocker aiozmq alembic pyzmq SQLAlchemy nvidia-docker	Backend.AI-manager Backend.AI-agent
	Go		Backend.AI-Jail
	ETC.	Docker 17.9	
Backend.AI <b>Cloud</b> CodeOnWeb.com	Python 3.6	Django 1.11	Customized Django Site-independent overloading
	Javascript ES6 / HTML5	Polymer 2.2	Lablup-webcomponents

# Backend.AI Structure (detail)



# Backend.AI Open-Source Project



- Communication
  - Slack → Synology Chat (+Synology CloudStation) → Microsoft Teams (+OneNote+OneDrive)
- Issue tracker
  - GitHub + Microsoft Teams
- Autodeploy
  - Lablup.AI : GitHub + post-commit signal to Azure → Autobuild on Azure webapp
  - CodeOnWeb: manual via script execution (to keep users safe)
- CI
  - Travis CI + Datadog
- Error reporting
  - Sentry + Datadog



This organization


backend.ai

Pull requests

Issues

Marketplace

Explore



Lablup

lab | up: Make AI Accessible - A start-up to innovate research / education processes.

📍 Seoul, Republic of Korea

🌐 <http://www.lablup.com>✉ [contact@lablup.com](mailto:contact@lablup.com)

📁 Repositories 68

👤 People 8

👥 Teams 4

📁 Projects 1

⚙ Settings

## Pinned repositories

Customize pinned repositories

≡ backend.ai

Meta-repository for backend.ai

Python ★ 34 🍴 5

≡ backend.ai-client-py

Client API Library for Lablup Backend.AI Service

Python ★ 4 🍴 1

≡ talkativot

Talkativot: Do-It-Yourself backbone for your AI friend

Python ★ 5 🍴 3

backend.ai

Type: All ▾

Language: All ▾

📁 New

16 results for repositories matching backend.ai

✕ Clear filter

backend.ai

Meta-repository for backend.ai

python api docker documentation distributed-computing

Python ★ 34 🍴 5 📄 LGPL-3.0 Updated 2 days ago

## Top languages

Python HTML JavaScript  
TypeScript C

## Most used topics

Manage

sorna python



## PACKAGE INDEX &gt;&gt;

[Browse packages](#)  
[List trove classifiers](#)  
[RSS \(latest 40 updates\)](#)  
[RSS \(newest 40 packages\)](#)  
[Terms of Service](#)  
[PyPI Tutorial](#)  
[PyPI Security](#)  
[PyPI Support](#)  
[PyPI Bug Reports](#)  
[PyPI Discussion](#)  
[PyPI Developer Info](#)

## ABOUT &gt;&gt;

## NEWS &gt;&gt;

## DOCUMENTATION &gt;&gt;

## DOWNLOAD &gt;&gt;

## COMMUNITY &gt;&gt;

## FOUNDATION &gt;&gt;

## CORE DEVELOPMENT &gt;&gt;

## backend.ai 1.0.2

## Lablup Backend.AI Meta-package

Downloads ↓

Backend.AI is a streamlined backend service framework hosting heterogeneous programming languages and popular AI frameworks. It manages the underlying computing resources for multi-tenant computation sessions where such sessions are spawned and executed instantly on demand.

All sub-projects are licensed under LGPLv3+.

By installing this meta-package, you get the client with command-line interface by default and optionally you may add the manager and agent using pip extra tags.

```
$ pip install backend.ai
(installs the common and client libs which includes CLI)
$ pip install backend.ai[manager]
(installs the common and client libs with the manager/gateway daemon)
$ pip install backend.ai[agent]
(installs the common and client libs with the agent daemon)
```



## Server-side Components

*Manager with API Gateway*

It routes external API requests from front-end services to individual agents. It also monitors and scales the cluster of multiple agents (a few tens to hundreds).

- Package namespace: `ai.backend.gateway` and `ai.backend.manager`
- <https://github.com/lablup/backend.ai-manager>

## Not Logged In

[Login](#)  
[Register](#)  
[Lost Login?](#)  
[Login with OpenID](#)   
[Login with Google](#) 

## Status

[Nothing to report](#)



## USER MANUALS

[API Overview](#)[Python Client Library](#)

## API COMMON REFERENCE

[API and Document Conventions](#)[Authentication](#)[Rate Limiting](#)[JSON Object References](#)

## USER API REFERENCE

[Introduction](#)[Kernel Management](#)[Code Execution \(Query Mode\)](#)[Code Execution and Monitoring  
\(Streaming Mode\)](#)[Code Execution \(Batch Mode\)](#)[Virtual Folders](#)

# Backend.AI Documentation

**Latest API version: v3.20170615 (beta)**

Backend.AI is a hassle-free backend for AI programming and service. It runs arbitrary user codes safely in resource-constrained environments, using Docker and our own sandbox wrapper.

Backend.AI supports various programming languages and runtimes, such as Python 2/3, R, PHP, C/C++, Java, Javascript, Julia, Octave, Haskell, Lua and NodeJS, as well as AI-oriented libraries such as TensorFlow, Keras, Caffe, and MXNet.

## FAQ

### vs. Notebooks

Product	Role	Problem and Solution
Apache Zeppelin, Databricks Notebook	Notebook-style document + code	Insecure host resource management



# TensorFlow Eager Execution





# TensorFlow: Summary



## ■ Statistics

- More than 24000 commits since Dec. 2015
- More than 1140 committers
- More than 6400 TensorFlow-related repository created on GitHub

## ■ Current

- Complete ML model prototyping
- Distributed training
- CPU / GPU / TPU / Mobile support

## ■ TensorFlow Serving

- Enables easier inference / model serving

## ■ XLA compiler (1.0~)

- Support various environments / speedups

## ■ Keras API Support (1.2~)

- High-level programming API
- Keras-compatible API

## ■ Eager Execution (1.4~)

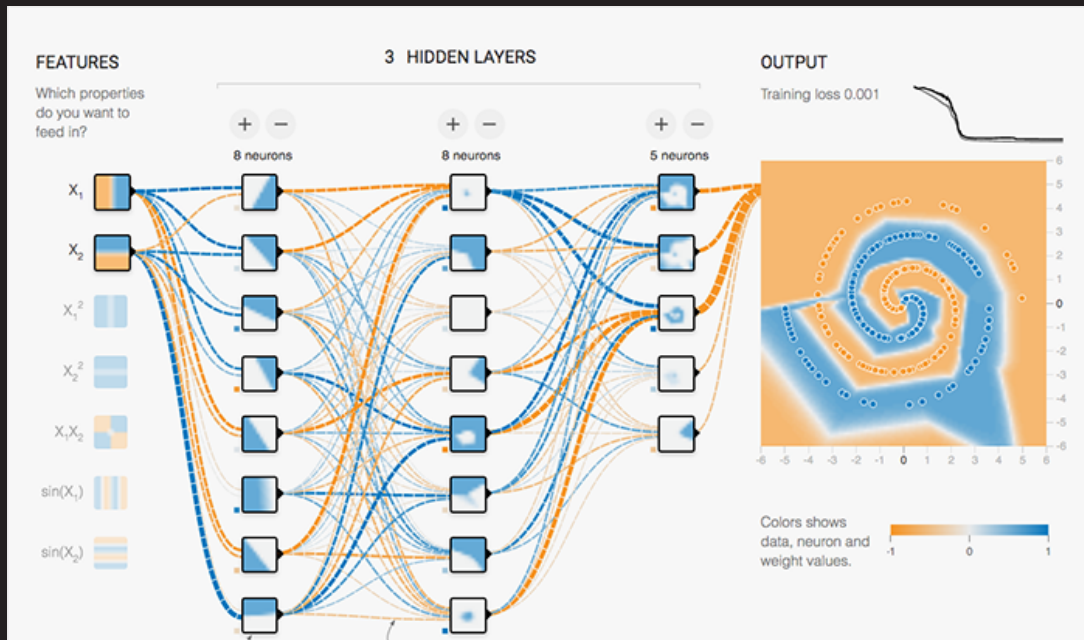
- Interactive mode of TensorFlow
- Treat TensorFlow python code as real python code



# How TensorFlow works



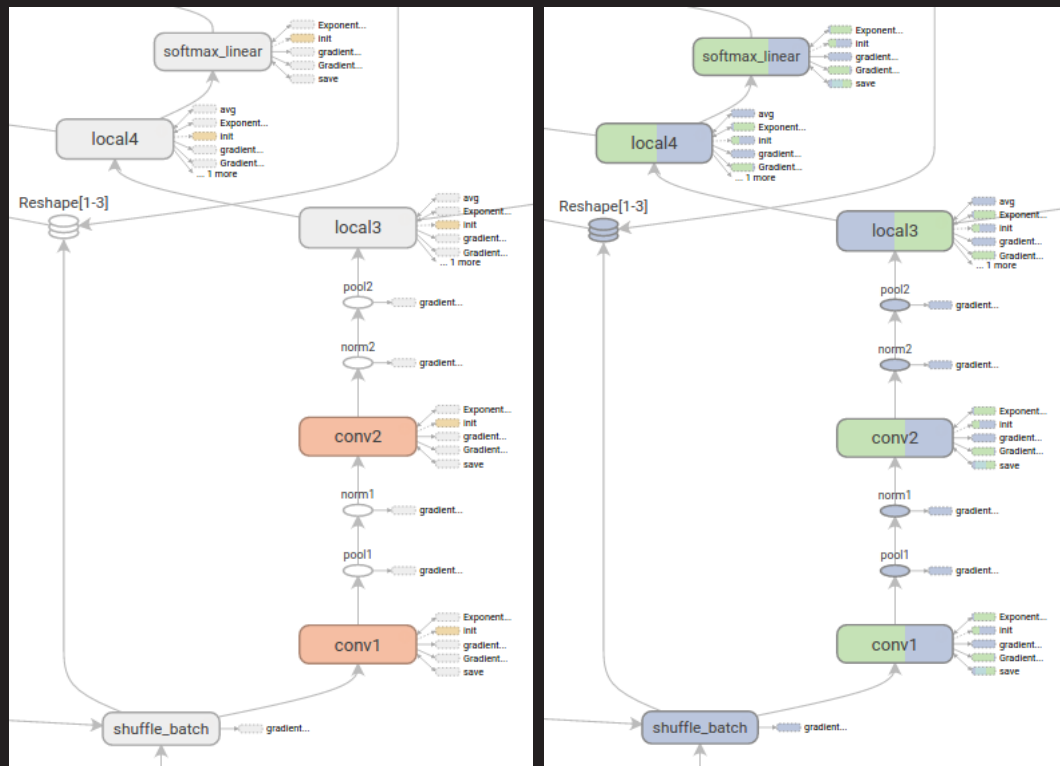
- CPU
  - Multiprocessor
    - ✓ AVX-based acceleration
    - ✓ GPU part in chip
  - OpenMP
- GPU
  - CUDA (NVidia) → cuDNN
  - OpenAL (AMD)
- TPU (1<sup>st</sup>, 2<sup>nd</sup> gen.)
  - ASIC for accelerating matrix calculation
  - In-house development by Google



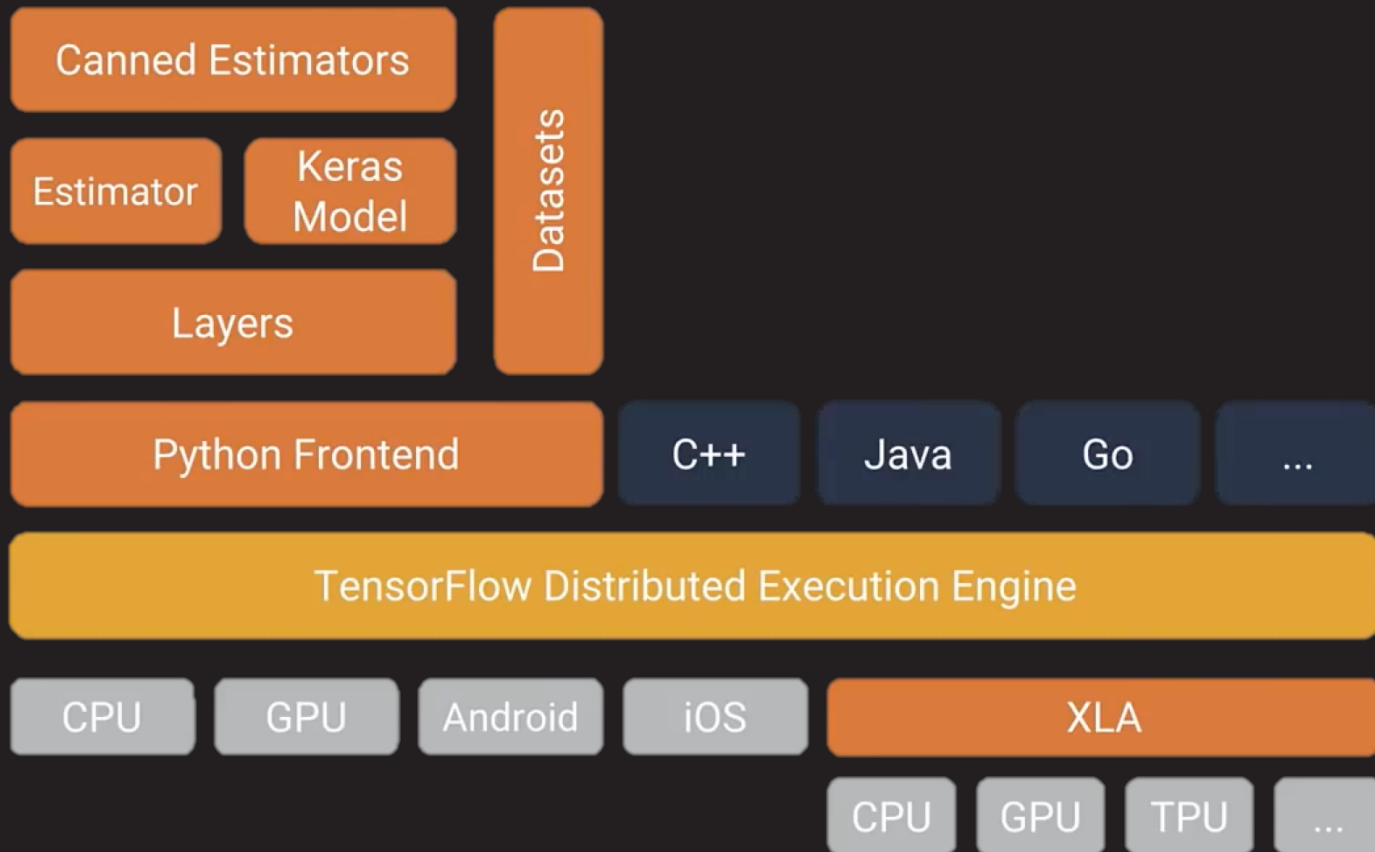
# How TensorFlow works



- Python but not Python
  - Python API is default API for TensorFlow
  - However, TF core is written in C++, with CuDNN library (for GPU acceleration)
- Computation Graph
  - User TF code is not a code
    - ✓ it is a **configuration** to generate **computation graph**
  - Session
    - ✓ Creates a computation graph and run the training using C++ core
  - Tedious debug process



# How TensorFlow works



# Eager execution



- Announced at Oct. 30, 2017
- Makes TensorFlow execute operations immediately
  - Returns concrete values
- Provides
  - A NumPy-like library for numerical computation
  - Support for GPU acceleration and automatic differentiation
  - A flexible platform for machine learning research and experiments
- Advantages
  - Python debugger tools
  - Immediate error reporting
  - Easy control flow
  - Python data structures

# Playing Eager Execution On Backend.AI



# Setting up Backend.AI

---



I

▲

# Eager execution: Session



```
x = tf.placeholder(tf.float32, shape=[1, 1])
m = tf.matmul(x, x)

print(m)
# Tensor("MatMul:0", shape=(1, 1), dtype=float32)

with tf.Session() as sess:
    m_out = sess.run(m, feed_dict={x: [[2.]]})
print(m_out)
# [[4.]]
```

```
x = [[2.]]
m = tf.matmul(x, x)

print(m)
# tf.Tensor([[4.]], dtype=float32, shape=(1, 1))
```



# Eager execution: Instant error



```
x = tf.gather([0, 1, 2], 7)
```

```
InvalidArgumentError: indices = 7 is not in [0, 3) [Op:Gather]
```

# Eager execution: removing metaprogramming



```
x = tf.random_uniform([2, 2])

with tf.Session() as sess:
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            print(sess.run(x[i, j]))
```

```
x = tf.random_uniform([2, 2])

for i in range(x.shape[0]):
    for j in range(x.shape[1]):
        print(x[i, j])
```

# Eager execution: Python Control Flow



```
a = tf.constant(6)
while not tf.equal(a, 1):
    if tf.equal(a % 2, 0):
        a = a / 2
    else:
        a = 3 * a + 1
print(a)
```

# Outputs

```
tf.Tensor(3, dtype=int32)
tf.Tensor(10, dtype=int32)
tf.Tensor(5, dtype=int32)
tf.Tensor(16, dtype=int32)
tf.Tensor(8, dtype=int32)
tf.Tensor(4, dtype=int32)
tf.Tensor(2, dtype=int32)
tf.Tensor(1, dtype=int32)
```

# Eager execution: Gradients



```
def square(x):  
    return tf.multiply(x, x)  # Or x * x  
  
grad = tfe.gradients_function(square)  
gradgrad = tfe.gradients_function(lambda x: grad(x)[0])  
  
print(square(3.))      # tf.Tensor(9., dtype=tf.float32)  
print(grad(3.))        # [tf.Tensor(6., dtype=tf.float32)]  
print(gradgrad(3.))    # [tf.Tensor(2., dtype=tf.float32)]
```

# Eager execution: Custom Gradients



```
def log1pexp(x):  
    return tf.log(1 + tf.exp(x))  
grad_log1pexp = tfe.gradients_function(log1pexp)  
  
print(grad_log1pexp(0.))
```

*Works fine, prints [0.5]*

# Eager execution: Custom Gradients



```
def log1pexp(x):  
    return tf.log(1 + tf.exp(x))  
grad_log1pexp = tfe.gradients_function(log1pexp)  
  
print(grad_log1pexp(100.))
```

*[nan] due to numeric instability*

# Eager execution: Custom Gradients



```
@tf.eager.custom_gradient
def log1pexp(x):
    e = tf.exp(x)
    def grad(dy):
        return dy * (1 - 1 / (1 + e))
    return tf.log(1 + e), grad
grad_log1pexp = tf.eager.gradients_function(log1pexp)

# Gradient at x = 0 works as before.
print(grad_log1pexp(0.))    # [0.5]
# And now gradient computation at x=100 works as well.
print(grad_log1pexp(100.))  # [1.0]
```

# Eager execution: Using GPUs



`tf.device()` for manual placement

```
with tf.device("/gpu:0"):
    x = tf.random_uniform([10, 10])
    y = tf.matmul(x, x)
    # x and y reside in GPU memory
```



# Eager execution: Building Models



The same APIs as graph building

(`tf.layers`, `tf.train.Optimizer`, `tf.data` etc.)

```
model = tf.layers.Dense(units=1, use_bias=True)
```

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1)
```

# Eager execution: Building Models



```
model = tf.layers.Dense(units=1, use_bias=True)
optimizer = tf.train.GradientDescentOptimizer(learning_rate=0.1)

# Define a loss function
def loss(x, y):
    return tf.reduce_mean(tf.square(y - model(x)))
```



## Compute and apply gradients

```
grad_fn = tfe.implicit_gradients(loss)

for (x, y) in get_next_batch():
    optimizer.apply_gradients(grad_fn(x, y))
```

# Status

---



- Alpha/Preview version out now!
- Single GPU, ResNet benchmark performance comparable to graphs
- Overheads on smaller operations is high
- Watch the release notes for upcoming TensorFlow releases for updates

# Backend.AI: How to join?



- <https://github.com/lablup/backend.ai>
  - We look forward to participating in Backend.AI development!
    - ✓ Would it be better if you come up with a pick? 🤔
  - V1.0 release! (Oct. 2017)
    - ✓ Installation and development manual provided! (Finally!)
- Future roadmap
  - Scheduler enhancements
    - ✓ Hybrid cloud - on-premise binding
  - Enhanced auto-scaling
    - ✓ Scale-in protection for long-time execution sessions
    - ✓ Cold/hot instance group management, depending on available cpu/memory/gpu slot capacity

# THANK YOU

If you have question, please contact via [contact@lablup.com](mailto:contact@lablup.com) !